



TMA 4300  
Computer Intensive Statistical Methods

## Exercise 2

*Gunnar Grung Grotmol*  
*Jan Henrik Jahren*  
*Audun Aass Engstrøm*

March 11, 2020  
Trondheim

FAKULTET FOR INFORMASJONSTEKNOLOGI OG  
ELEKTRONIKK

INSTITUTT FOR MATEMATISKE FAG

# Project 2

## Problem A

In this problem we look at the coal mining data, and do an analysis using Markov chain Monte Carlo. The plot of accumulated disasters from the year 1851 to 1962 is given below.

```
library(coda)
library(MASS)
library(boot)
library(MCMCpack)
```

```
## ##
## ## Markov Chain Monte Carlo Package (MCMCpack)

## ## Copyright (C) 2003-2020 Andrew D. Martin, Kevin M. Quinn, and Jong Hee Park

## ##
## ## Support provided by the U.S. National Science Foundation

## ## (Grants SES-0350646 and SES-0350613)
## ##
```

```
library(ggplot2)
library(dae)
library("INLA")
```

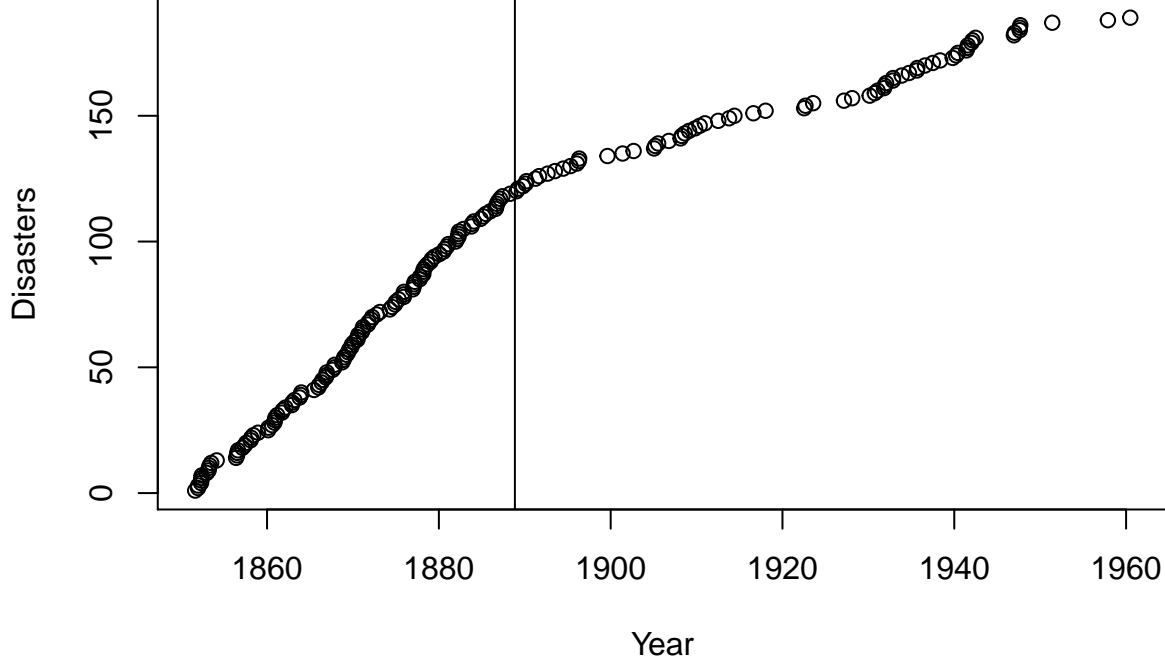
```
## Loading required package: Matrix
```

```
## Loading required package: sp
```

```
## Loading required package: parallel
```

```
## This is INLA_19.09.03 built 2019-09-03 09:07:31 UTC.
## See www.r-inla.org/contact-us for how to get help.
## To enable PARDISO sparse library; see inla.pardiso()
```

```
par(mfrow=c(1,1))
x.all <- coal[[1]]
x <- x.all[2:(length(x.all)-1)]
# Plot coalmine-data
plot(x,1:length(x),ylab = "Disasters",xlab = "Year"); abline(v = mean(x))
```



As seen from the above plot, the number of incidents seem to be greater in the 19'th century compared to the 20'th. The slope to the left of the mean(horisontal line) is greater than to the right. The coal mining disasters are assumed to follow an inhomogenius Poisson process with intensity function  $\lambda(t)$ , which is assumed to be piece-wise constant with  $n$  segments. We let  $t_0$  and  $t_{n+1}$  denote the start en end time of the data set. It is possible to show that the likelihood function for the observed data is then

$$f(x|t_1, \dots, t_n, \lambda_1, \dots, \lambda_n) = \exp\left(-\int_{t_0}^{t_{n+1}} \lambda(t)dt\right) \prod_{k=0}^n \lambda_k^{y_k} = \exp\left(-\sum_{k=0}^n \lambda_k(t_{k+1} - t_k)\right) \prod_{k=0}^n \lambda_k^{y_k}$$

Where  $y_k$  is the number of incidents during the period  $t_k$  to  $t_{k+1}$ . We assume that all  $t_i$ , for  $i \in \{1, \dots, n\}$ , are independent and uniformly distributed between  $t_0$  and  $t_{n+1}$ . Apriori, we assume all  $\lambda_i$  to be independent of  $t_i$  and  $\lambda_j$  for all  $i, j \in \{1, \dots, n\}$  and  $j \neq i$ . Furthermore, we assume all  $\lambda_i$  to be i.i.d. from the same gamma distribution with shape parameter  $\alpha = 2$  and scale parameter  $\beta$ . In other words

$$f(\lambda_i|\beta) = \frac{1}{\beta^2} \lambda_i e^{-\frac{\lambda_i}{\beta}}.$$

Moreover, the scale hyperparameter is assigned the improper prior

$$f(\beta) \propto \frac{e^{-1/\beta}}{\beta} \quad \text{for } \beta > 0.$$

We now restrict ourselves to look at  $n = 1$ , and therefore the model parameters  $\theta = (t_1, \lambda_0, \lambda_1, \beta)$ . The posterior distribution can be found (up to a normalizing constant)

$$\begin{aligned} f(\theta|x) &\propto f(x|\theta)f(\theta) = f(x|\theta)f(t_1)f(\lambda_0, \lambda_1|\beta)f(\beta) \\ &= f(x|\theta)f(t_1)f(\lambda_0|\beta)f(\lambda_1|\beta)f(\beta) \\ &\propto \frac{\lambda_0^{y_0+1} \lambda_1^{y_1+1}}{\beta^5} e^{-\lambda_0(t_1-t_0)-\lambda_1(t_2-t_1)-\frac{1+\lambda_0+\lambda_1}{\beta}}. \end{aligned}$$

Furthermore the full conditionals can be derived. The F.C. for  $t_1$  is

$$\begin{aligned} f(t_1|x, \lambda_0, \lambda_1, \beta) &\propto f(t_1)f(x|\lambda_0, \lambda_1, t_1, \beta) \\ &\propto \lambda_0^{y_0} \lambda_1^{y_1} e^{-\lambda_0(t_1-t_0)-\lambda_1(t_2-t_1)} \propto \lambda_0^{y_0} \lambda_1^{y_1} e^{t_1(\lambda_1-\lambda_0)} \end{aligned}$$

which is not a known distribution. Next, the full conditional  $\lambda_0$  is given by

$$\begin{aligned} f(\lambda_0 \mid \lambda_1, t_1, \beta, x) &\propto f(\lambda_0 \mid \beta) f(\lambda_1, t_1, x \mid \lambda_0, \beta) \\ &\propto f(\lambda_0 \mid \beta) f(x \mid \lambda_1, \lambda_0, t_1, \beta) \\ &\propto \lambda_0^{y_0+1} e^{-\lambda_0(t_1-t_0)-\lambda_0/\beta} \end{aligned}$$

which is the gamma distribution,  $\lambda_0 \mid \lambda_1, t_1, \beta, x \sim \text{Gamma}(y_0 + 2, t_1 - t_0 + 1/\beta)$  using shape and rate parameters. Using the exact same reasoning, the F.C. for  $\lambda_1$  is

$$f(\lambda_1 \mid \lambda_0, t_1, \beta, x) \propto \lambda_1^{y_1+1} e^{-\lambda_1(t_2-t_1)-\lambda_1/\beta}.$$

Therefore,  $\lambda_1 \mid \lambda_0, t_1, \beta, x \sim \text{Gamma}(y_1 + 2, t_2 - t_1 + 1/\beta)$  using shape and rate parameters. Lastly, the full conditional of  $\beta$  is

$$\begin{aligned} f(\beta \mid \lambda_0, \lambda_1, t_1, x) &\propto f(\beta) f(\lambda_0, \lambda_1, t_1, x \mid \beta) \propto f(\beta) f(\lambda_0 \mid \beta) f(\lambda_1 \mid \beta) \\ &\propto \frac{1}{\beta^5} e^{-\beta(1+\lambda_0+\lambda_1)}, \end{aligned}$$

which is the inverse gamma distribution. That is,  $\beta \mid \lambda_0, \lambda_1, t_1, x \sim \text{InvGamma}(4, (1 + \lambda_0 + \lambda_1))$  with shape and scale parameter respectively.

Next, we want to sample from the posterior distribution using a single site MCMC. The variable of interest when using Metropolis-Hastings is  $t_1$ . The other variables will be drawn from their known full conditionals. The new values are found in the order:  $t_1$ ,  $\beta$ ,  $\lambda_0$  and lastly  $\lambda_1$ . We want the limiting distribution of the Metropolis-Hastings step to be

$$f(\theta \mid x) \propto f(t_1 \mid x, \lambda_0, \lambda_1, \beta) \propto \lambda_0^{y_0} \lambda_1^{y_1} e^{-\lambda_0(t_1-t_0)-\lambda_1(t_2-t_1)}.$$

keeping the other variables constant. Therefore, we introduce the notation  $\pi(t) = f(t \mid x, \lambda_0, \lambda_1, \beta)$ . Let the probability for transitioning from state  $t_1$  to  $t'_1$  to be  $p(t'_1 \mid t_1)$ . In order to assure a unique limiting distribution we need the Markov chain to be aperiodic and reversible. We require

$$\pi(t_1)p(t'_1 \mid t_1) = \pi(t'_1)p(t_1 \mid t'_1),$$

otherwise, the Markov chain would not be reversible, as for example a transition from  $t_1$  to  $t'_1$  would occur too often. Since this might not be the case, we should only accept a transition from  $t_1$  to  $t'_1$  with a probability  $\alpha(t_1, t'_1)$ . We need to construct  $\alpha$  such that

$$\pi(t_1)q(t'_1 \mid t_1)\alpha(t_1, t'_1) = \pi(t'_1)q(t_1 \mid t'_1)\alpha(t'_1, t_1),$$

where  $p(t'_1 \mid t_1) = q(t'_1 \mid t_1)\alpha(t_1, t'_1)$ . Here,  $q(t'_1 \mid t)$  is the distribution used to propose a new candidate  $t'_1$ . A natural choice for  $\alpha$  is to pick

$$\alpha(t_1, t'_1) = \min\left\{1, \frac{\pi(t'_1)q(t_1 \mid t'_1)}{\pi(t_1)q(t'_1 \mid t_1)}\right\}.$$

In our implementation, we choose the random walk Markov chain, with the symmetric transition probability  $q(t'_1 \mid t_1) \sim \mathcal{N}(t_1, \sigma^2)$  for some properly chosen tuning parameter  $\sigma$ . Since the transition probability is symmetric,  $q(t'_1 \mid t_1) = q(t_1 \mid t'_1)$ , the acceptance probability becomes

$$\alpha(t_1, t'_1) = \min\left\{1, \frac{\pi(t'_1)}{\pi(t_1)}\right\} = \min\left\{1, \frac{\lambda_0^{y'_0} \lambda_1^{y'_1} e^{-\lambda_0 t'_1 + \lambda_1 t'_1}}{\lambda_0^{y_0} \lambda_1^{y_1} e^{-\lambda_0 t_1 + \lambda_1 t_1}}\right\}.$$

In order to implement the algorithm, we introduce the variables and help functions

```

#global variables
t0 <- x.all[1]
t2 <- x.all[length(x.all)]
x <- x.all[2:(length(x.all)-1)]
#help functions
get.y <- function(x, t) {
  # Get y_0 and y_1, the number of events occuring before and after t
  y_0 <- sum(x <= t)
  y_1 <- length(x) - y_0
  return(c(y_0,y_1))
}
sample.FC.beta <- function(lambda_0,lambda_1) {
  # Sample beta from inverse gamma distribution
  return(rinvgamma(1,shape = 4, scale = (lambda_0 + lambda_1 + 1)))
}
sample.FC.lambda_0 <- function(t_1, beta, x) {
  # Sample lambda_0 from gamma distribution
  return(rgamma(1,shape = get.y(x,t_1)[1] + 2,scale = beta/(beta*(t_1-t0)+1)))
}
sample.FC.lambda_1 <- function(t_1, beta,x) {
  # Sample lambda_1 from gamma distribution
  lambda_1 = rgamma(1,shape = get.y(x,t_1)[2] + 2, scale = beta/(beta*(t2-t_1)+1))
}

```

Since the values of  $\lambda_0^{y_0} \lambda_1^{y_1}$  can be very large the calculations are done in logarithmic scale in order to increase stability. The code for the single site Markov chain Monte Carlo is given below.

```

single_site_MH_MCMC <- function(N,t_1,lambda_0, lambda_1, beta, x, sigma) {
  #create data frame and insert first values
  MCMC.data <- data.frame(matrix(0,N,4))
  colnames(MCMC.data) <- c('t_1','lambda_0','lambda_1','beta')
  MCMC.data <- data.frame(t_1 = double(), lambda_0 = double(), lambda_1 = double(), beta = double())
  MCMC.data[1,] <- c(t_1,lambda_0,lambda_1,beta)

  successes <- 0 # count number of sucessfull new draws
  for (i in 2:N){
    t.new <- rnorm(1,mean = t_1, sd = sigma)#propose new t_1 value
    if( (t.new < t0) | (t.new > t2)){
      # if outside allowed region, throw away new sample
      log_alpha <- -Inf
    }
    else{
      # Find y-values at t_1 and t.new
      y <- get.y(x, t_1)
      y.new <- get.y(x, t.new)

      #calculate numerator and denominator of acceptance ratio log scale

      numerator <- y.new[1]*log(lambda_0) + y.new[2]*log(lambda_1) + t.new*(lambda_1 - lambda_0)

      denominator <- y[1]*log(lambda_0) + y[2]*log(lambda_1) + t_1*(lambda_1 - lambda_0)
    }
  }
}

```

```

    log_alpha <- numerator - denominator      # log Acceptance probability
  }
  # if u is U(0,1), u < alpha, accept new draw,
  if (runif(1) < exp(log_alpha)) {
    t_1 <- t.new
    successes <- successes + 1 # count number of successes
  }
  # if u > alpha, do nothing
  # update data frame
  MCMC.data[i,] <- c(t_1,lambda_0,lambda_1,beta)
  #draw new lambda_0, lambda_1, beta
  beta <- sample.FC.beta(lambda_0,lambda_1)
  lambda_0 <- sample.FC.lambda_0(t_1,beta,x)
  lambda_1 <- sample.FC.lambda_1(t_1,beta,x)
  # do everything again with new parameters
}
return(c(MCMC.data,successes/N))
}

```

We next try the algorithm to check burn in and mixing properties. We set the initial value of  $t_1$  to the mean of  $t_0$  and  $t_2$ . For convenience we set the initial value of the hyperparameter  $\beta = 1$  and sample  $\lambda_0$  and  $\lambda_1$  from the their full conditionals above as starting values.

```

# initial values of single site algorithm
t.init = 0.5*(t0+t2)
beta = 1.0
lambda_0 = sample.FC.lambda_0(t.init, beta, x)
lambda_1 = sample.FC.lambda_1(t.init, beta, x)
lambda_0;lambda_1;

```

```
## [1] 2.795728
```

```
## [1] 0.8669295
```

We run for 5000 iterations with different values of  $\sigma$ . Three values are compared. These are  $\sigma = 1$ ,  $\sigma = 5$  and  $\sigma = 10$ .

```

# steps of Markov chain
N <- 10000
# with sigma = 1
sigma = 1
results <- single_site_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)
MCMC.data1 = results[1:4]; success_rate1 = results[5]

# with sigma = 5
sigma = 5
results <- single_site_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)
MCMC.data2 = results[1:4]; success_rate2 = results[5]

# with sigma = 10
sigma = 10
results <- single_site_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)

```

```
MCMC.data3 = results[1:4]; success_rate3 = results[5]

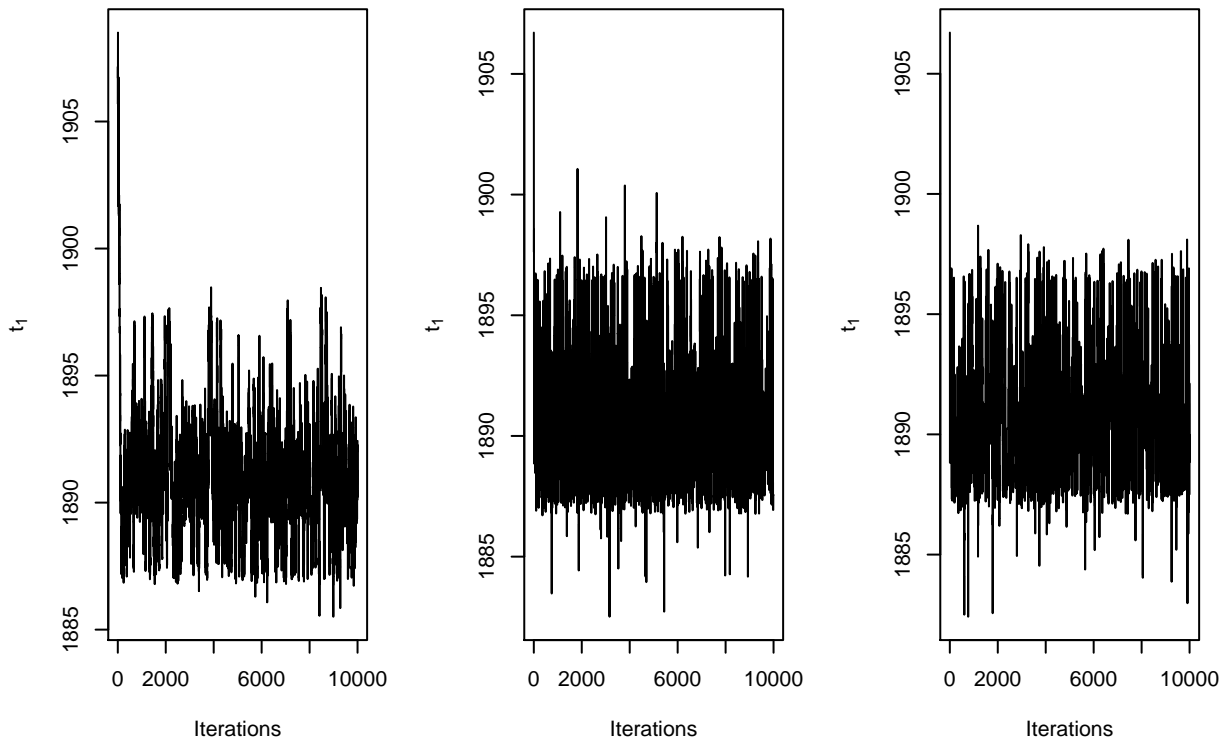
success_rate1; success_rate2; success_rate3;
```

```
## [[1]]
## [1] 0.6064
```

```
## [[1]]
## [1] 0.3218
```

```
## [[1]]
## [1] 0.1783
```

```
# Plot results
par(mfrow=c(1,3))
plot(MCMC.data1$t_1, type='l', ylab=expression(t[1]), xlab = "Iterations")
plot(MCMC.data2$t_1, type='l', ylab=expression(t[1]), xlab = "Iterations")
plot(MCMC.data3$t_1, type='l', ylab=expression(t[1]), xlab = "Iterations")
```

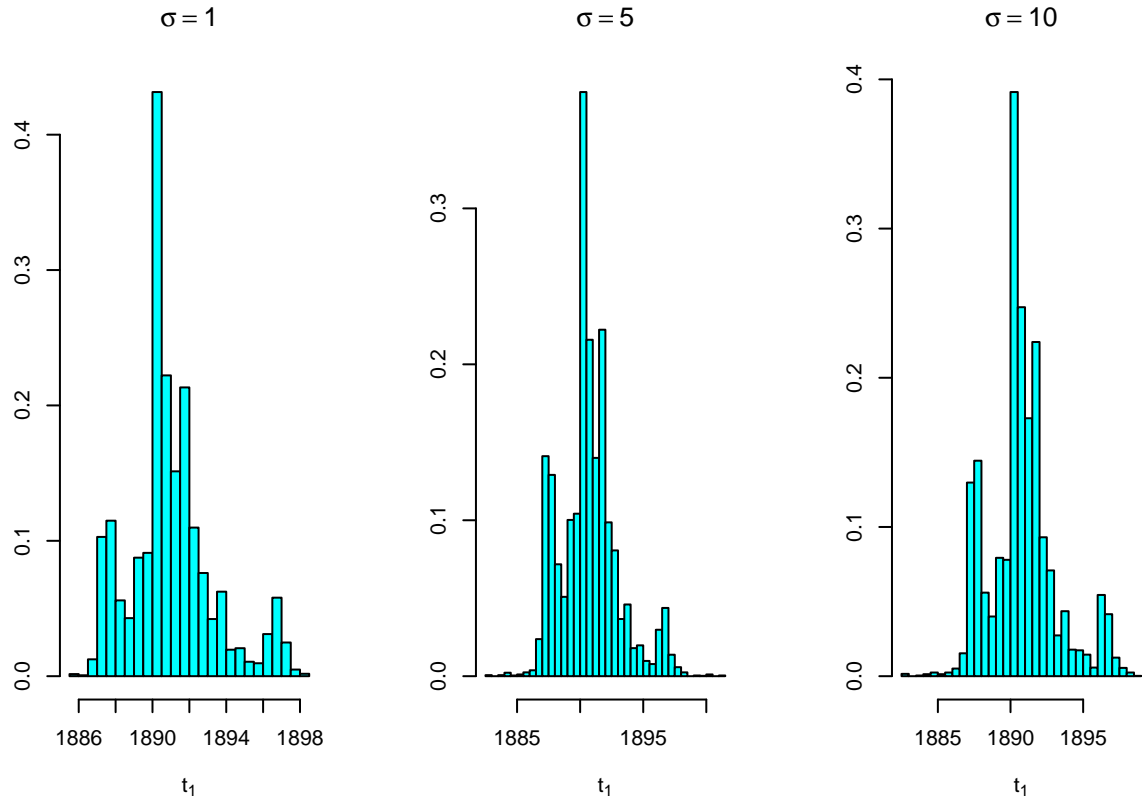


As seen from the above plot, the algorithm seems to use somewhere close to 500 iterations to reach the points around the equilibrium state for  $\sigma = 1$ . For both  $\sigma = 5$  and  $\sigma = 10$ , the sampled values converge to the equilibrium faster. The acceptance rates for the different values of  $\sigma$  are quite different, with the acceptance ratio for  $\sigma = 1$  around 0.6,  $\sigma = 5$  around 0.3 and  $\sigma = 10$  around 0.2. If the value of sigma is set too high, many proposed values will be too extreme, yielding a lower convergence rate. The higher values of sigma in other words introduce smaller acceptance values. However, we see that all Markov chains oscillate around the same line around 1890. In other words, the algorithm does converge to the same distribution for all  $\sigma$  values, but at different rates. This is seen in the plot below, with a burn in of 1000.

```

Burn.in = 1000
par(mfrow=c(1,3))
truehist(MCMC.data1$t_1[Burn.in:N], xlab = expression(t[1]),main = expression(sigma == 1))
truehist(MCMC.data2$t_1[Burn.in:N], xlab = expression(t[1]),main = expression(sigma == 5))
truehist(MCMC.data3$t_1[Burn.in:N], xlab = expression(t[1]),main = expression(sigma ==10))

```



The above plot shows that convergence is achieved using any  $\sigma$ . Furthermore, the limiting distribution is unique. Since  $\sigma = 1$  uses around 500 iterations in order to reach an equilibrium, the burn in will be set to be a reasonable 1000. Using this value, we run the algorithm further using  $\sigma = 5$ , and compare the results to the Coal mining disaster data.

```

#Parameters
N <- 50000
sigma <- 5
Burn.in <- 1000
MCMC.data <- single_site_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)[0:4]

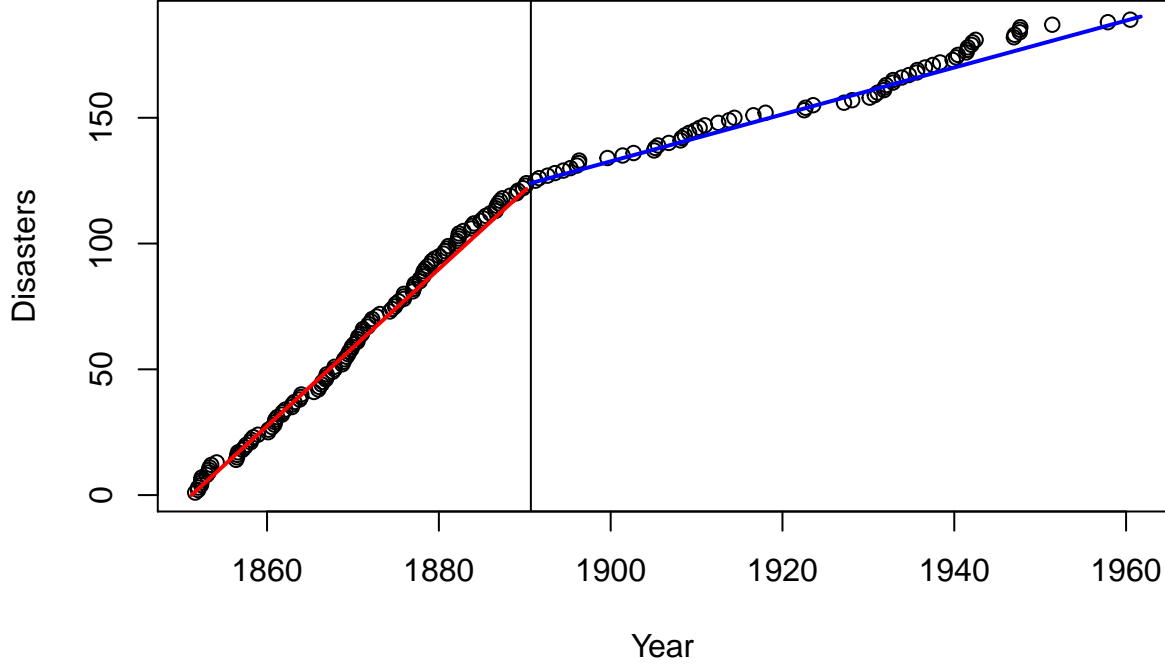
```

```

# see if algorithm makes sense
t_avg = mean(MCMC.data$t_1[Burn.in:N])
lambda_0_avg = mean(MCMC.data$lambda_0[Burn.in:N])
lambda_1_avg = mean(MCMC.data$lambda_1[Burn.in:N])
#plot
par(mfrow=c(1,1))
plot(x,1:length(x),xlab = "Year", ylab = "Disasters")
abline(v = mean(MCMC.data$t_1), ylab = "Disasters")
lines(seq(t0, t_avg), lambda_0_avg*(seq(t0, t_avg)-t0), type = "l", col = "red", lwd=2)
lines(seq(t_avg, t2), lambda_1_avg*(seq(t_avg, t2)-t_avg)+get.y(x,t_avg)[1],type = "l", col = "blue", lwd=2)

```





```
# Average values of the rate parameters
sprintf("mean of t_1: %.3f,", t_avg)
```

```
## [1] "mean of t_1: 1890.703,"
```

```
sprintf("mean of lambda_0: %.3f,", lambda_0_avg)
```

```
## [1] "mean of lambda_0: 3.120,"
```

```
sprintf("mean of lambda_1: %.3f ", lambda_1_avg)
```

```
## [1] "mean of lambda_1: 0.933 "
```

The above plot shows a vertical line with x-value equal to the mean of sampled  $t_1$ . Furthermore, it shows a line with slope equal to the sample mean of  $\lambda_0$  to the left, and a line with slope mean of  $\lambda_1$  to the right. These lines closely hug the cumulative distribution of the disaster data. Therefore, the sampled values make sense. Our algorithm is most likely implemented correctly.

We next define a two block Metropolis-Hastings algorithm when sampling from  $f(\theta|x)$ . The first block proposal is updating the block  $(t_1, \lambda_0, \lambda_1)$  keeping  $\beta$  constant. In this block we generate a proposal to a new value for  $t_1$ ,  $\tilde{t}_1$  using a normal distribution  $\tilde{t}_1 \sim \mathcal{N}(t_1, \sigma_{t_1}^2)$ , with mean  $t_1$  and variance  $\sigma_{t_1}^2$ . Thereafter, we propose the new values of  $(\lambda_0, \lambda_1)$ ,  $(\tilde{\lambda}_0, \tilde{\lambda}_1)$ , using their joint full conditional  $f(\lambda_0, \lambda_1|x, \tilde{t}_1, \beta)$  inserted the value  $\tilde{t}_1$ . The second block is  $(\lambda_0, \lambda_1, \beta)$  keeping  $t_1$  unchanged. A new proposal to  $\beta$ ,  $\tilde{\beta}$ , is sampled from the normal distribution  $\tilde{\beta} \sim \mathcal{N}(\beta, \sigma_\beta^2)$ , with mean  $\beta$ . The proposed values  $(\tilde{\lambda}_0, \tilde{\lambda}_1)$  are thereafter found using their full conditional  $f(\lambda_0, \lambda_1|x, t_1, \tilde{\beta})$  inserted  $\tilde{\beta}$ . Let  $\alpha_1$  and  $\alpha_2$  be the acceptance probability of the first and second block respectively. Focusing on the first block, we again require

$$\alpha_1(t_1, \lambda_0, \lambda_1, \tilde{t}_1, \tilde{\lambda}_0, \tilde{\lambda}_1) = \min \left\{ 1, \frac{f(\tilde{\theta}|x)q_1(t_1, \lambda_0, \lambda_1|\tilde{t}_1, \tilde{\lambda}_0, \tilde{\lambda}_1)}{f(\theta|x)q_1(\tilde{t}_1, \tilde{\lambda}_0, \tilde{\lambda}_1|t_1, \lambda_0, \lambda_1)} \right\}$$

Where

$$q_1(\tilde{t}_1, \tilde{\lambda}_0, \tilde{\lambda}_1 | t_1, \lambda_0, \lambda_1, x, \beta) = f(\tilde{t}_1 | t_1, \lambda_0, \lambda_1, x, \beta) f(\tilde{\lambda}_0, \tilde{\lambda}_1 | \tilde{t}_1, t_1, \lambda_0, \lambda_1, x, \beta) = f(\tilde{t}_1 | t_1) f(\lambda_0 | x, \tilde{t}_1, \beta) f(\lambda_1 | x, \tilde{t}_1, \beta).$$

Since the value of  $\tilde{t}_1$  is sampled from the symmetric normal distribution with mean  $t_1$ , we have  $f(\tilde{t}_1 | t_1) = f(t_1 | \tilde{t}_1)$ . Therefore, this part of the expression cancels out. Furthermore, the exponential part and the polynomial part of the gamma distribution cancels with the posterior distribution. Therefore the acceptance probability for the block 1 transition becomes

$$\alpha_1(t_1, \lambda_0, \lambda_1, \tilde{t}_1, \tilde{\lambda}_0, \tilde{\lambda}_1) = \min \left\{ 1, \frac{\Gamma(\tilde{y}_0 + 2) \Gamma(\tilde{y}_1 + 2)}{\Gamma(y_0 + 2) \Gamma(y_1 + 2)} \frac{(t_1 - t_0 + 1/\beta)^{y_0+2} (t_2 - t_1 + 1/\beta)^{y_1+2}}{(\tilde{t}_1 - t_0 + 1/\beta)^{\tilde{y}_0+2} (t_2 - \tilde{t}_1 + 1/\beta)^{\tilde{y}_1+2}} \right\}.$$

we now look at the transition probability of block 2. The transition kernel,  $q_2$ , is on the form

$$q_2(\tilde{\lambda}_0, \tilde{\lambda}_1, \tilde{\beta} | \lambda_0, \lambda_1, \beta) = f(\tilde{\beta} | \lambda_0, \lambda_1, \beta) f(\tilde{\lambda}_0, \tilde{\lambda}_1 | \tilde{\beta}, \lambda_0, \lambda_1, \beta) = f(\tilde{\beta} | \beta) f(\tilde{\lambda}_0, \tilde{\lambda}_1 | x, t_1, \tilde{\beta})$$

yielding

$$\alpha_2(\tilde{\lambda}_0, \tilde{\lambda}_1, \tilde{\beta}, \lambda_0, \lambda_1, \beta) = \min \left\{ 1, \frac{\beta^5}{\tilde{\beta}^5} e^{1/\beta - 1/\tilde{\beta}} \frac{(t_1 - t_0 + 1/\beta)^{y_0+2} (t_2 - t_1 + 1/\beta)^{y_1+2}}{(t_1 - t_0 + 1/\tilde{\beta})^{y_0+2} (t_2 - t_1 + 1/\tilde{\beta})^{y_1+2}} \right\}.$$

The implementation can be seen below. When implementing we set  $\sigma_{t_1} = 2\sigma_\beta = \sigma$ .

```
Block_MH_MCMC <- function(N,t_1, lambda_0, lambda_1, beta, x, sigma){
  #create data frame and insert first values
  MCMC.data <- data.frame(matrix(0,N,4))
  colnames(MCMC.data) <- c('t_1','lambda_0','lambda_1','beta')
  MCMC.data <- data.frame(t_1 = double(), lambda_0 = double(), lambda_1 = double(), beta = double())
  MCMC.data[1,] <- c(t_1,lambda_0,lambda_1,beta)

  for (i in 2:N){
    #block 1
    t.new <- rnorm(1,mean=t_1,sd = sigma)
    #draw new sample if outside allowed region
    while(t.new < t0 | t.new > t2){t.new <- rnorm(1,mean=t_1,sd = sigma)}
    lambda_0.new <- sample.FC.lambda_0(t.new,beta, x)
    lambda_1.new <- sample.FC.lambda_1(t.new,beta, x)
    y <- get.y(x,t_1)
    y.new <- get.y(x,t.new)

    numerator <- lfactorial(y.new[1]+2) + lfactorial(y.new[2]+2) +(y[1]+2)*log(t_1-t0+1/beta) +

    denominator <- lfactorial(y[1]+2) + lfactorial(y[2]+2) +(y.new[1]+2)*log(t.new - t0+1/beta) + (y.new[2]+2)*log(t_2-t.new+1/beta)
    alpha_log <- numerator - denominator
    if (runif(1)< exp(alpha_log)){
      t_1 <- t.new
      lambda_0 <- lambda_0.new
      lambda_1 <- lambda_1.new
    }

    #block 2
    beta.new <- rnorm(1,mean= beta,sd = 0.5*sigma)
```

```

# draw new sample if beta non-positive
while(beta.new <= 0){beta.new <- rnorm(1,mean= beta,sd = 0.5*sigma)}
lambda_0.new <- sample.FC.lambda_0(t_1,beta.new, x)
lambda_1.new <- sample.FC.lambda_1(t_1,beta.new, x)
y <- get.y(x,t_1)

numerator <- 5*log(beta)+(1/beta-1/beta.new) +(y[1]+2)*log(t_1-t0+1/beta) + (y[2]+2)*log(t2-t_1+
denominator <- 5*log(beta.new) +(y[1]+2)*log(t_1-t0+1/beta.new)+ (y[2]+2)*log(t2-t_1+1/beta.new)

alpha_log <- numerator - denominator
if (runif(1)< exp(alpha_log)){
  lambda_0 <- lambda_0.new
  lambda_1 <- lambda_1.new
  beta <- beta.new
}
MCMC.data[i,] <- c(t_1,lambda_0,lambda_1,beta)
}
return(MCMC.data)
}

```

We run the algorithm for different choices of  $\sigma$  to evaluate the mixing properties and burn in. The values picked are  $\sigma \in \{1, 5, 10, 20\}$ .

```

# initial values of single site algorithm
t.init <- 0.5*(t0+t2)
beta <- 1.0
lambda_0 <- sample.FC.lambda_0(t.init, beta, x)
lambda_1 <- sample.FC.lambda_1(t.init, beta, x)
# steps of Markov chain
N <- 10000
# with sigma = 1
sigma = 1
MCMC.data1 <- Block_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)
# with sigma = 5
sigma = 5
MCMC.data2 <- Block_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)
# with sigma = 10
sigma = 10
MCMC.data3 <- Block_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)
# with sigma = 20
sigma = 20
MCMC.data4 <- Block_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)

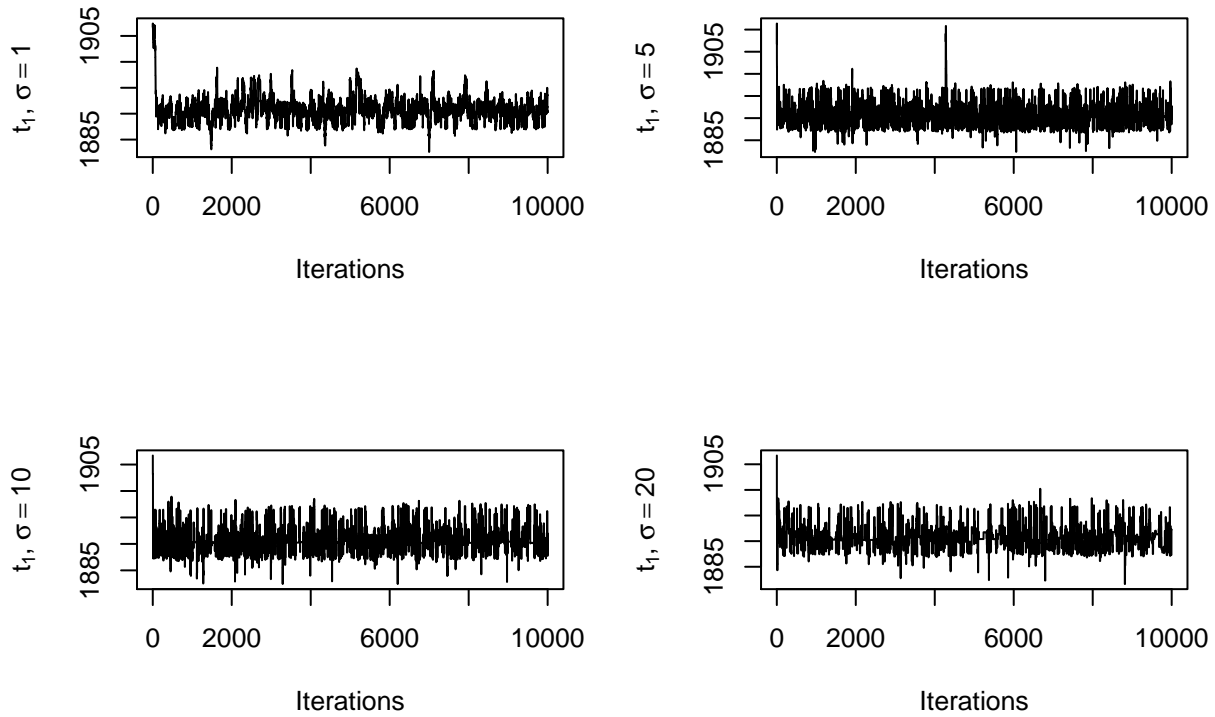
```

Below are two plots showing the sampled values of  $t_1$  and  $\beta$ .

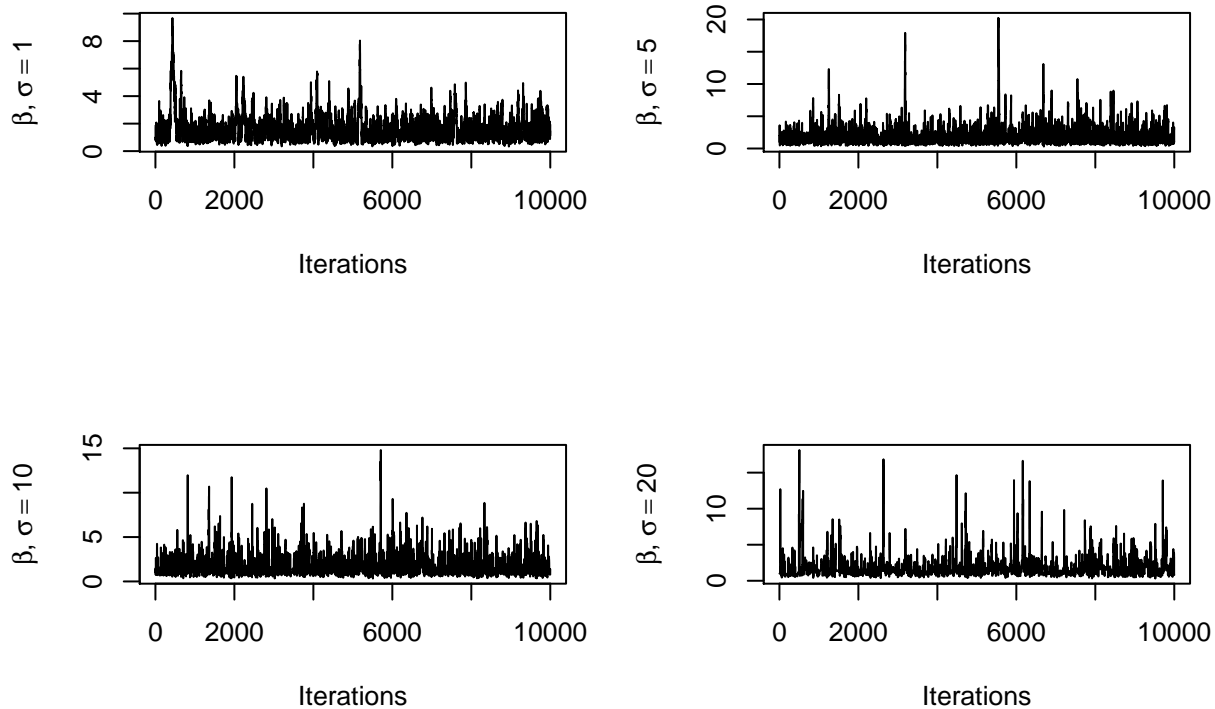
```

# plot t_1
par(mfrow=c(2,2))
plot(MCMC.data1$t_1, type='l', ylab=expression(paste(t[1], " ", sigma== 1)), xlab = "Iterations")
plot(MCMC.data2$t_1, type='l', ylab=expression(paste(t[1], " ", sigma== 5)), xlab = "Iterations")
plot(MCMC.data3$t_1, type='l', ylab=expression(paste(t[1], " ", sigma==10)), xlab = "Iterations")
plot(MCMC.data4$t_1, type='l', ylab=expression(paste(t[1], " ", sigma==20)), xlab = "Iterations")

```



```
#plot beta
par(mfrow=c(2,2))
plot(MCMC.data1$beta, type='l', ylab=expression(paste(beta, " ", sigma== 1)), xlab = "Iterations")
plot(MCMC.data2$beta, type='l', ylab=expression(paste(beta, " ", sigma== 5)), xlab = "Iterations")
plot(MCMC.data3$beta, type='l', ylab=expression(paste(beta, " ", sigma==10)), xlab = "Iterations")
plot(MCMC.data4$beta, type='l', ylab=expression(paste(beta, " ", sigma==20)), xlab = "Iterations")
```

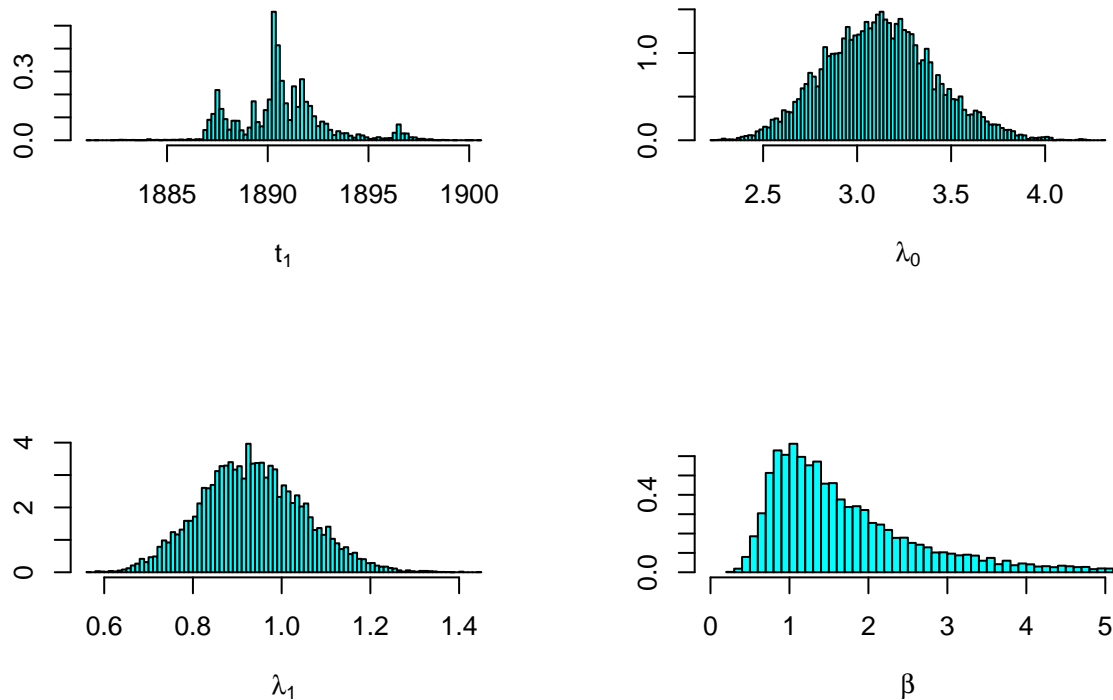


As seen from the above plots, the optimal mixing occurs for  $\sigma = 5$ . For  $\sigma = 1$  the acceptance value is high with small movements, leading to slow convergence. For high values of  $\sigma$ ,  $\sigma = 10$  and  $\sigma = 20$  the mixing is

slower than for  $\sigma = 5$ . This is especially clear for  $\sigma = 20$ , where the line is less noisy, especially for  $t_1$ . This implies a low acceptance rate. It is, however, evident that the mixing property of the two block Markov chain Monte Carlo is greater than it is for the single site MCMC. The number of iterations needed to reach equilibrium for the two block seems far less than 500, as was the value for the single site MCMC. We pick  $\sigma = 5$  and run the algorithm further. We then apply the a burn in of 1000, and look at the results.

```
N <- 50000
Burn.in <- 1000
MCMC.data <- Block_MH_MCMC(N,t.init,lambda_0,lambda_1,beta,x,sigma)[Burn.in:N,]
```

```
par(mfrow=c(2,2))
truehist(MCMC.data$t_1, xlab = expression(t[1]))
truehist(MCMC.data$lambda_0, xlab = expression(lambda[0]))
truehist(MCMC.data$lambda_1, xlab = expression(lambda[1]))
truehist(MCMC.data$beta, xlim = c(0,5),, xlab = expression(beta))
```



```
sprintf("Mean of t_1: %.2f",mean(MCMC.data$t_1))
```

```
## [1] "Mean of t_1: 1890.66"
```

```
sprintf("Mean of lambda_0: %.3f",mean(MCMC.data$lambda_0))
```

```
## [1] "Mean of lambda_0: 3.120"
```

```
sprintf("Mean of lambda_1: %.3f",mean(MCMC.data$lambda_1))
```

```
## [1] "Mean of lambda_1: 0.934"
```

```
sprintf("Mean of beta: %.3f",mean(MCMC.data$beta))
```

```
## [1] "Mean of beta: 1.854"
```

```
sprintf("Covariance lambda_0, lambda_1 : %.5f",cov(MCMC.data$lambda_0,MCMC.data$lambda_1))
```

```
## [1] "Covariance lambda_0, lambda_1 : 0.00120"
```

Above are the estimated distributions of  $f(t_1|x)$ ,  $f(\lambda_0|x)$ ,  $f(\lambda_1|x)$  and  $f(\beta|x)$ . In addition, the estimated values of  $E[t_1|x]$ ,  $E[\lambda_0|x]$ ,  $E[\lambda_1|x]$ ,  $E[\beta|x]$  and  $\text{Cov}[\lambda_0, \lambda_1|x]$

## Problem B

1)

```
#Read in from the website and plot the data
```

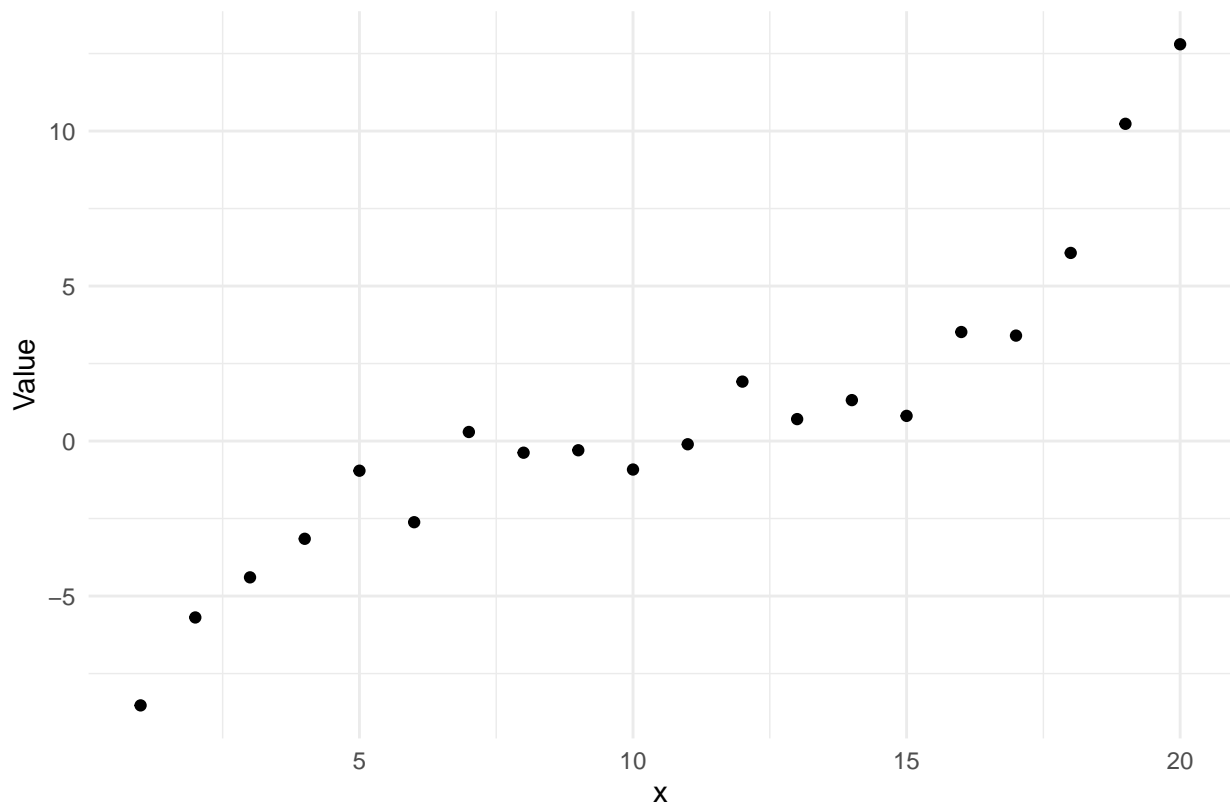
```
filepath <- "https://www.math.ntnu.no/emner/TMA4300/2020v/Exercise/exercise2/Gaussiandata.txt"
```

```
gaus<- read.table(file = filepath, header = FALSE)
```

```
x <- seq(1, length(gaus$V1))
```

```
gaus <- cbind(x, gaus)
```

```
plot <- ggplot(data = gaus, aes(x = x, y = V1)) + geom_point() + theme_minimal() + labs(y = "Value")  
plot + labs(caption = "Plot of the observed data")
```



Plot of the observed data

Our model can be written as a combination of the likelihood of the response and the prior distributuon of the latent field (Hierarchical Mode):

$$\mathbf{y} \mid \mathbf{f} \sim \prod_{t=1}^T P(y_t \mid \eta_t)$$

$$\mathbf{f} \mid \theta \sim \pi(\mathbf{f} \mid \theta) = N(\mathbf{0}, \mathbf{Q}(\theta)^{-1})$$

Where

$$\theta \sim \text{Gamma}(1, 1)$$

Since we can write our model in this way its by definition a latent Gaussian model.

Since our model is a LGM, each of our observations  $y_i$  only depends on one of the of the elements (linear predictors) in the latent Gaussian field and the dimension of hyperparameters is only 1. (Must be  $\leq 9$ ) . It is also desirable to use INLA in this case cause the cause the precision matrix  $\mathbf{Q}(\theta)$  is sparse.

## 2)

We want to implement a block Gibbs sampling using two proposed value for  $\theta$  and the linear predictor vector  $\eta$  from

$$\pi(\theta \mid \eta, y)$$

and

$$\pi(\eta \mid \theta, y)$$

respectively.

First we look at the posterior

$$\pi(\boldsymbol{\eta}, \theta \mid \mathbf{y}) \propto \pi(\theta) \pi(\boldsymbol{\eta} \mid \theta) \prod_{t=1}^T \pi(y_t \mid \eta_t, \theta)$$

Where

$$\pi(\theta) = \exp(-\theta)$$

(pdf of gamma(1,1))

$$\pi(\eta \mid \theta)$$

is the second order walk model

$$\prod_{t=1}^T \pi(y_t \mid \eta_t, \theta)$$

is the likelihood of the normal distribution. Taking the log likelihood and inserting the known paramteres:

$$\pi(\boldsymbol{\eta}, \theta \mid \mathbf{y}) \propto \frac{\theta^{(T-2)/2}}{(2\pi)^{1/2}} \exp(-\theta - \frac{\theta}{2} \sum_{t=3}^T (\eta_t - 2\eta_{t-1} + \eta_{t-2})^2 - \frac{1}{2} \sum_{t=1}^T (y_t - \eta_t)^2)$$

The full conditional for theta is then (removed the terms that doesnt depend on theta)

$$\pi(\theta \mid \mathbf{y}, \boldsymbol{\eta}) \propto \theta^{\frac{T}{2}-1} \exp(-\theta(1 + \frac{1}{2} \sum_{t=3}^T (\eta_t - 2\eta_{t-1} + \eta_{t-2})^2)$$

We recognize this as the gamma function with paramter  $\alpha = T/2$  and  $\beta = 1 + \frac{1}{2} \sum_{t=3}^T (\eta_t - 2\eta_{t-1} + \eta_{t-2})^2$ . By introducing  $\mathbf{Q}$  as the precision matrix to account for the “markov term” in the linear predictors we can write the term in matrix notation. We define this as

$$\mathbf{Q} = \theta \mathbf{L} \mathbf{L}^T$$

where  $\mathbf{L}$  is a  $T \times (T-2)$  matrix that accounts for the coefficients in the “markov term”.

$$\mathbf{L} = \begin{bmatrix} 1 & -2 & 1 & 0 & 0 & 0 & \dots \\ 0 & 1 & -2 & 1 & 0 & 0 & \dots \\ \vdots & & \ddots & \ddots & \ddots & & \\ 0 & 0 & 0 & 0 & 1 & -2 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & -2 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}.$$

The full conditional in matrix notation is then:

$$\pi(\theta \mid \mathbf{y}, \boldsymbol{\eta}) \propto \theta^{\frac{T}{2}-1} \exp(-\theta - \frac{1}{2} \boldsymbol{\eta}^T \mathbf{Q} \boldsymbol{\eta})$$

To get the exact gamma distribution we define  $\mathbf{Q}^* = \frac{1}{\theta} \mathbf{Q}$  and we end up with

$$\pi(\theta \mid \mathbf{y}, \boldsymbol{\eta}) \propto \theta^{\frac{T}{2}-1} \exp(-\theta(1 + \frac{1}{2} \boldsymbol{\eta}^T \mathbf{Q}^* \boldsymbol{\eta}))$$

Now for the full conditional for  $\boldsymbol{\eta}$

$$\pi(\boldsymbol{\eta} \mid \beta, \mathbf{y}) \propto \exp(-\frac{\theta}{2} \sum_{t=3}^T (\eta_t - 2\eta_{t-1} + \eta_{t-2})^2 - \frac{1}{2} \sum_{t=1}^T (y_t - \eta_t)^2)$$

In matrix notation

$$= \exp(-\frac{1}{2} (\boldsymbol{\eta}^T \mathbf{Q} \boldsymbol{\eta}) + (\mathbf{y} - \boldsymbol{\eta})^T (\mathbf{y} - \boldsymbol{\eta}))$$

$$= \exp(-\frac{1}{2} \boldsymbol{\eta}^T (\mathbf{Q} + \mathbf{I}) \boldsymbol{\eta} + \mathbf{y}^T \boldsymbol{\eta})$$

2)

```
#Code for the Precision matrix. Initialising theta to be 1 so it can generate Q and Q*
precMat <- function(Tdim, theta=1){
  L <- diag(1, Tdim, Tdim-2)
  L[row(L) - col(L) == 1] <- -2
  L[row(L) - col(L) == 2] <- 1
  Q <- theta*L%*%t(L)
  return(Q)
}
```



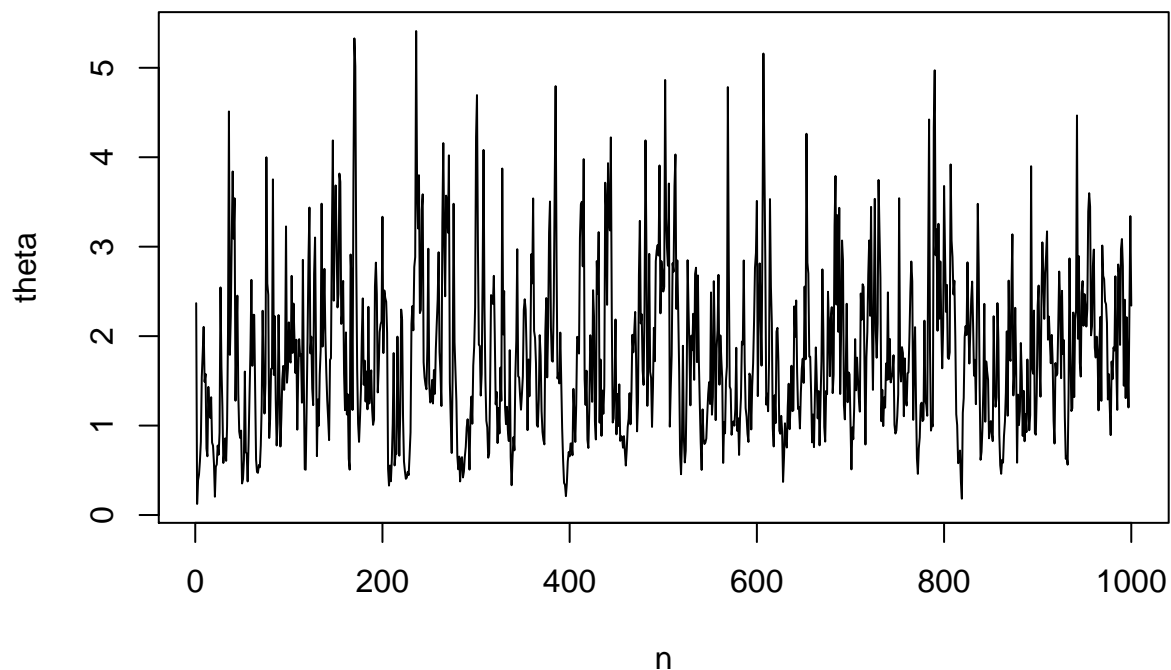
```

#Code for Gibbs sampling routine
gibbsSampling <- function(y, n){
  theta <- rep(0,n)
  Tdim <- length(y)
  A <- precMat(Tdim)
  Idmat <- diag(1,Tdim)
  eta <- matrix(0,n,Tdim)
  theta[1] <- rgamma(1,1,1)
  eta[1,] <- y

  #Iterating and calculate the next theta and eta based on the previous and conditional to each other
  for(i in 2:n){
    rate <- 1 + 0.5*t(eta[i-1,]) %*% A %*% eta[i-1,]
    theta[i] <- rgamma(1,shape = Tdim/2, rate = rate)
    var <- solve(A*theta[i] + Idmat , sparse = TRUE)
    mean <- var %*% y
    eta[i,] <- rmvnorm(mean, var)

  }
  samples <- cbind(theta, eta)
  return(samples)
}
test <- gibbsSampling(gaus$V1, 1000)
#Plot of the theta and first eta.
plot(test[,1], type="l",xlab="n", ylab="theta")

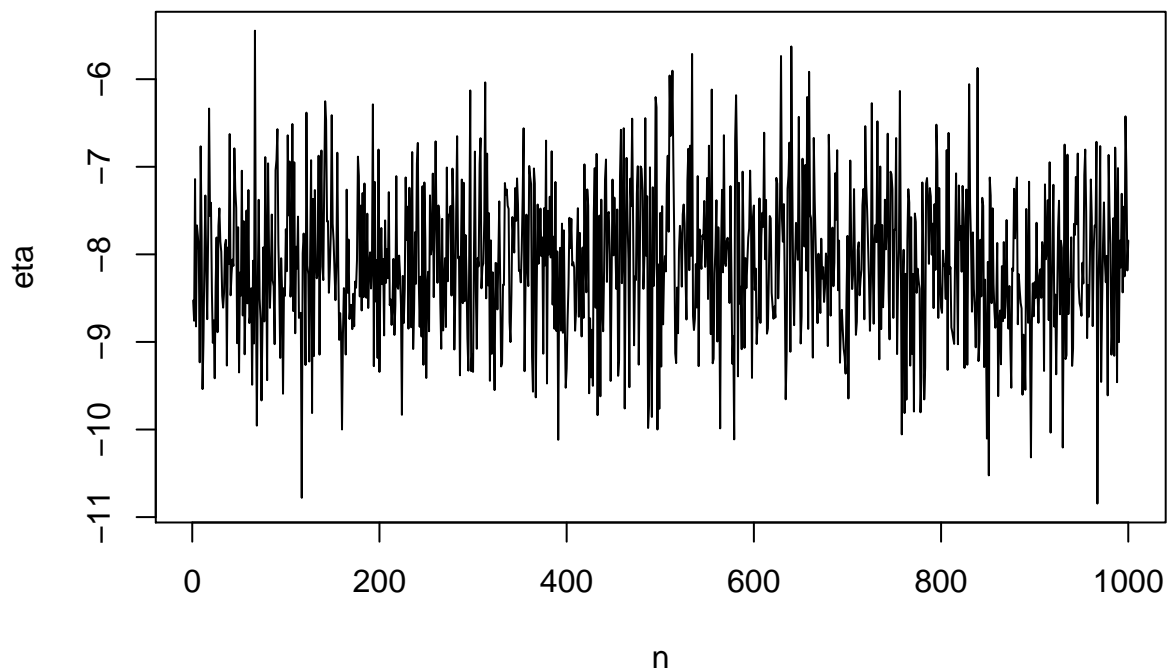
```



```

plot(test[,2], type="l",xlab="n", ylab="eta")

```



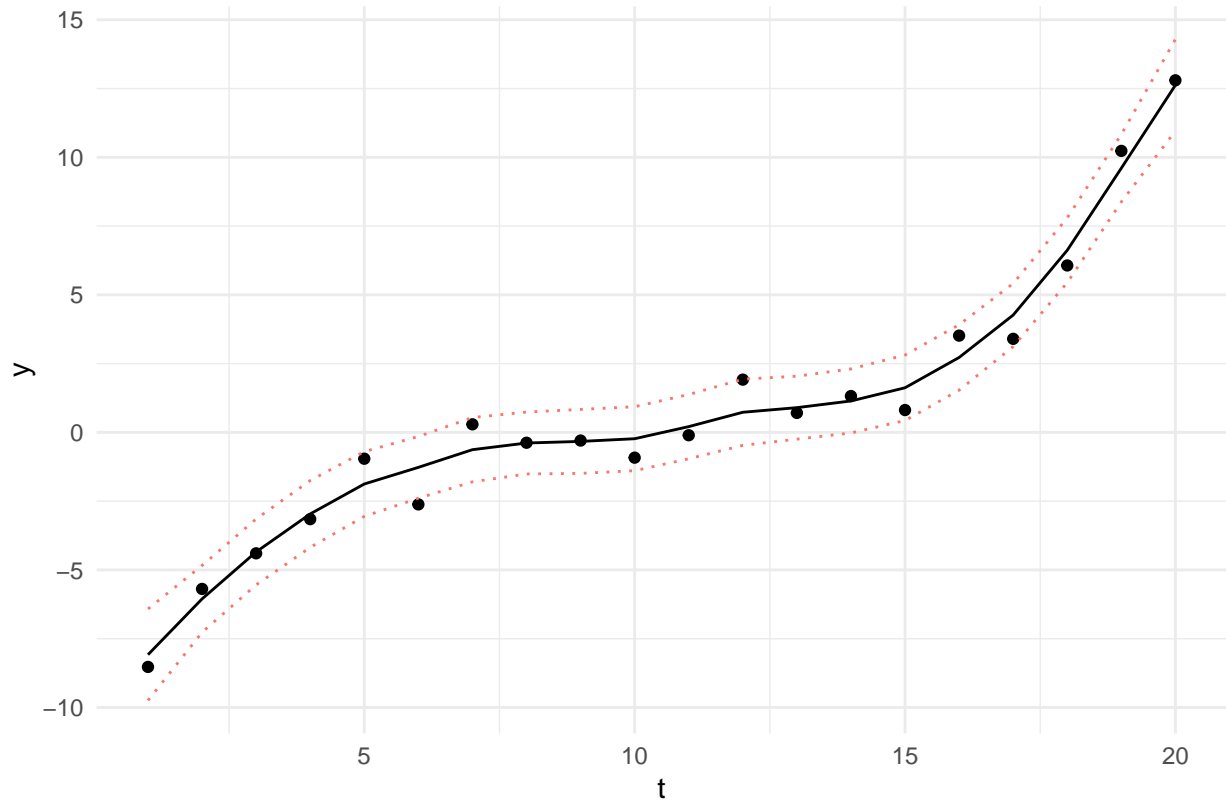
From these plots we don't notice any burn-in period for neither eta or theta. This is because we don't see any patterns that stand out for the first  $x$  iterations. However, to account for small patterns that are hard to notice, we choose to omit the first 100 iterations.

```
#Omitting the first 100 samples to get rid of the "burnn" period
newtest <- test[-c(1:100),]

meanvec <- colMeans(newtest)[-1]
varvec <- apply(newtest, 2, sd)[-1]
#Confidence interval
upperlim <- meanvec + qnorm(0.025, lower.tail = FALSE)*varvec
lowerlim <- meanvec - qnorm(0.025, lower.tail = FALSE)*varvec

test3 <- cbind(gaus, lowerlim, upperlim)

#Plot it
plot <- ggplot(data = test3, aes(x = x, y = meanvec)) + geom_point(aes(x = x, y = V1 )) + geom_line() +
plot <- plot + labs(y = "y", x = "t", caption = "Estimate of the smooth effect, where the dotted lines
plot + theme(legend.position="none")
```



Estimate of the smooth effect, where the dotted lines represents a 95% confidence interval

3)

$$\begin{aligned}\pi(\theta \mid \mathbf{y}) &\propto \frac{\pi(\mathbf{y} \mid \boldsymbol{\eta}, \theta) \pi(\boldsymbol{\eta} \mid \theta) \pi(\theta)}{\pi(\boldsymbol{\eta} \mid \theta, \mathbf{y})} \\ &\propto \frac{\theta^{\frac{T-2}{2}} \exp(-\theta - \frac{1}{2}(\boldsymbol{\eta}^T(\mathbf{Q} + \mathbf{I})\boldsymbol{\eta} + \mathbf{y}^T\mathbf{y} + \mathbf{y}^T)\boldsymbol{\eta})}{|\mathbf{Q} + \mathbf{I}| \exp(-\frac{1}{2}(\boldsymbol{\eta} - (\mathbf{Q} + \mathbf{I})^{-1}\mathbf{y})^T(\mathbf{Q} + \mathbf{I})(\boldsymbol{\eta} - (\mathbf{Q} + \mathbf{I})^{-1}\mathbf{y}))}\end{aligned}$$

The etas cancels out and we get

$$\propto \frac{\theta^{\frac{T-2}{2}}}{|\mathbf{Q} + \mathbf{I}|^{\frac{1}{2}}} \exp(-\theta - \frac{1}{2}\mathbf{y}^T(\mathbf{I} - (\mathbf{Q} + \mathbf{I})^{-1})\mathbf{y})$$

*#Function to calculate the pi of theta given y*

```
piTheta<- function(theta){
  Tdim <- length(gaus$V1)
  thet1 <- length(theta)
  pivec <- rep(0,thet1)
  det <- rep(0,thet1)
  Q <- precMat(Tdim)
  Idmat <- diag(Tdim)

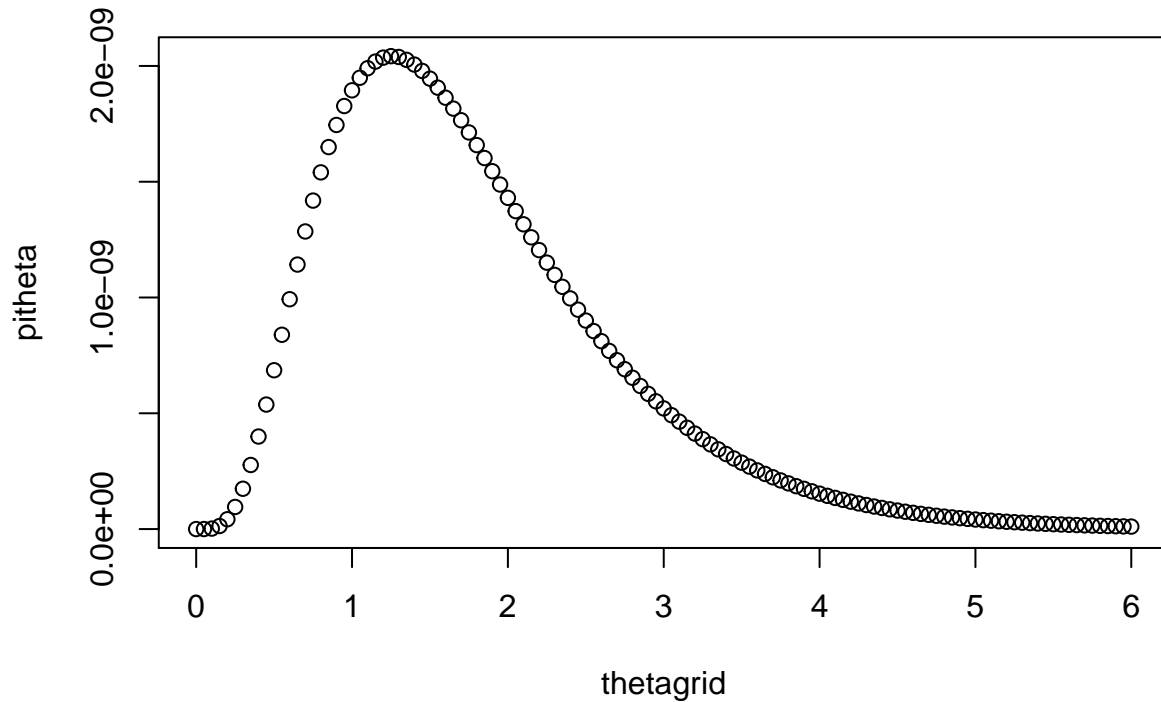
  for(i in 1:thet1){
    det <- det(Q*theta[i]+ Idmat)
```

```

    pivec[i] <- theta[i]^(Tdim/2 -1)/sqrt(det)* exp(-theta[i]-0.5*t(gaus$V1) %*% (Idmat-solve(Q*theta[i]
  }
  return(pivec)
}

#Defining the grid
delta <- 0.05
thetagrid <- seq(0,6,delta)
#simple plot
pitheta <- piTheta(thetagrid)
plot(thetagrid, pitheta)

```



The

plot shows the the values of  $\pi(\theta | y)$  for values of theta from 0-6.

Comparing with the MCMC estimate it looks pretty similar as we see the highest frequency of theta is between 0.5 and 2.5 in both plots. (Note that this graph is not normalized)

#### 4)

By making a grid and sum up we can make a approcimation of the integral (numerical integration) as follows:

$$\pi(\eta_i | y) = \int \pi(\eta_i | y, \theta) \pi(\theta | y) d\theta \approx \sum_{j=1}^J \pi(\eta_i | y, \theta^j) \pi(\theta^j | y) \Delta j$$

Where  $\Delta j$  is the steplength for the grid we made for theta in the previous task.

```

#Function to calculate pi of eta given y and theta for eta10

piEtaYTheta <- function(etagrid, theta, i = 10){
  Tdim <- length(gaus$V1)

```

```

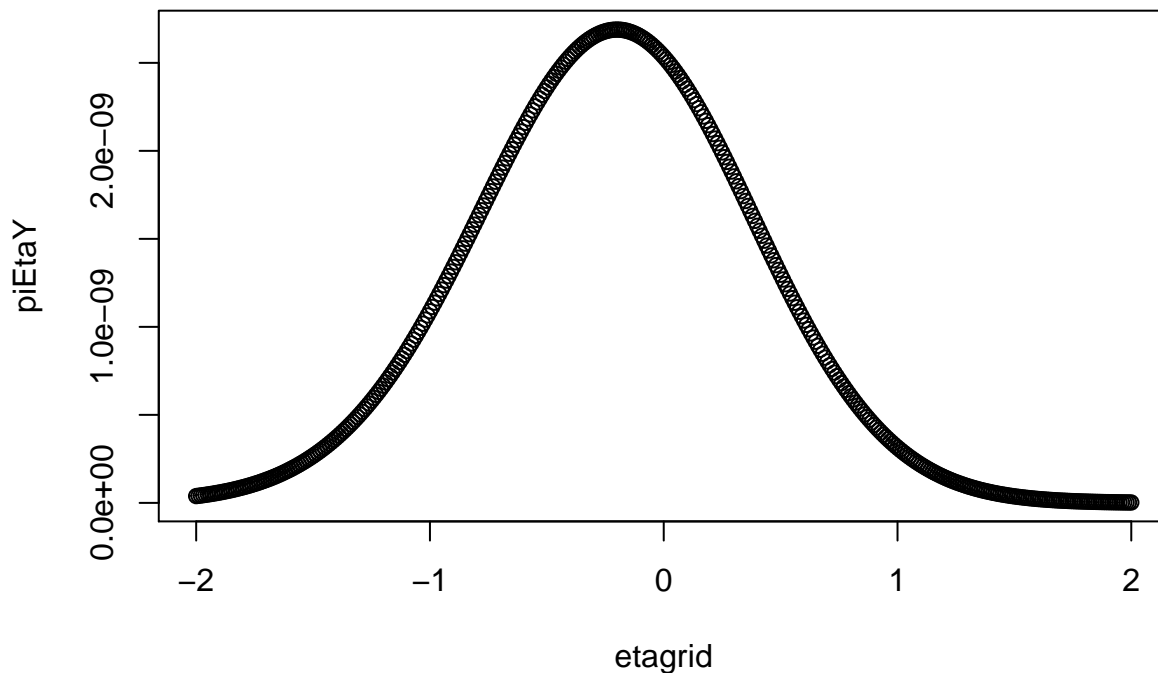
Q <- precMat(Tdim, theta)
A <- solve(Q + diag(Tdim))
mean <- (A%% gaus$V1)[i]
var <- A[i,i]
pi <- dnorm(etagrid, mean, sqrt(var))
return(pi)
}

#Approximation of the integral for i = 10 using numerical integration with the previous functions
piEtaY <- function(thetagrid, etagrid){
  sum <- rep(0, length(etagrid))
  steplength <- thetagrid[2]-thetagrid[1]
  thetay <- piTheta(thetagrid)
  for (j in 1:length(thetagrid)){
    theta <- thetagrid[j]
    sum <- sum + piEtaYTheta(etagrid, theta) * thetay[j]*steplength
  }
  return(sum)
}

#Constructing the etagrid
etagrid <- seq(-2,2,0.01)
piEtaY<- piEtaY(thetagrid, etagrid)

plot(etagrid, piEtaY)

```

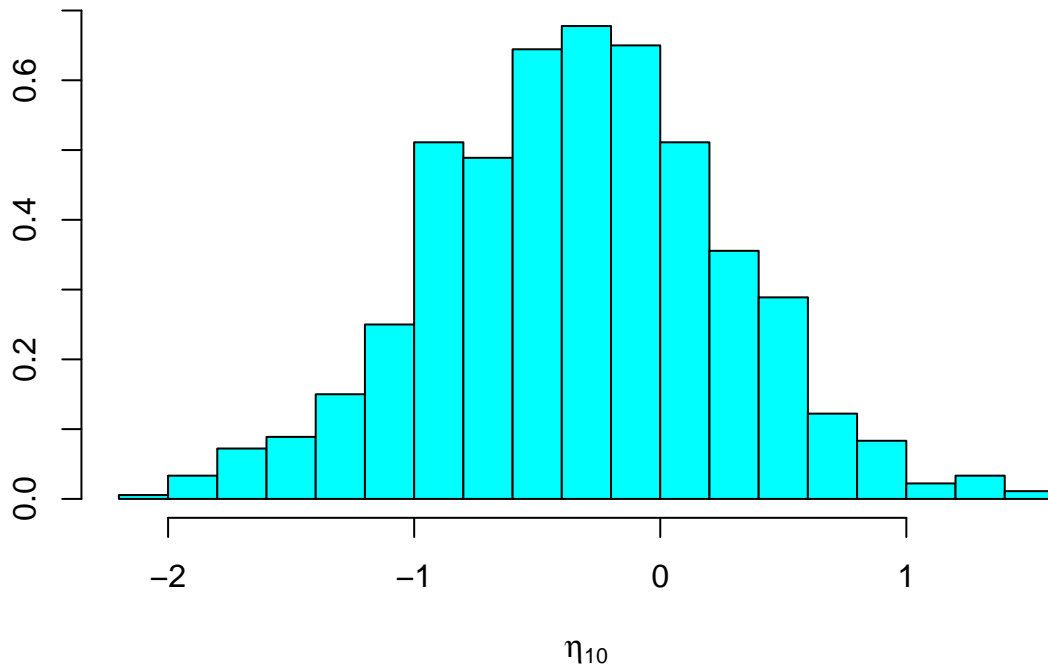


```

est2 <- newtest[,10]

#Simple plot
truehist(est2, xlab = bquote(eta[10]))

```



The first plot is the plot from INLA method and the second from the Gibbs sampling we obtained in task 2. Both of them show the marginal posterior distribution for the smooth effect for  $\eta_{10}$ . As we can see the mean is shifted slightly to the left of zero for both and they look quite similar.

5)

```
y <- gaus$V1
x <- gaus$x

#Using the built in INLA
thetapri <- list(theta = list(prior = "log.gamma", param = c(1,1)))

formula <- y ~ f(x, model = "rw2", hyper = thetapri, constr = FALSE) - 1

fit <- inla(formula = formula, family = "gaussian", data = gaus, control.family = list(hyper=list(prec = 1)))
```

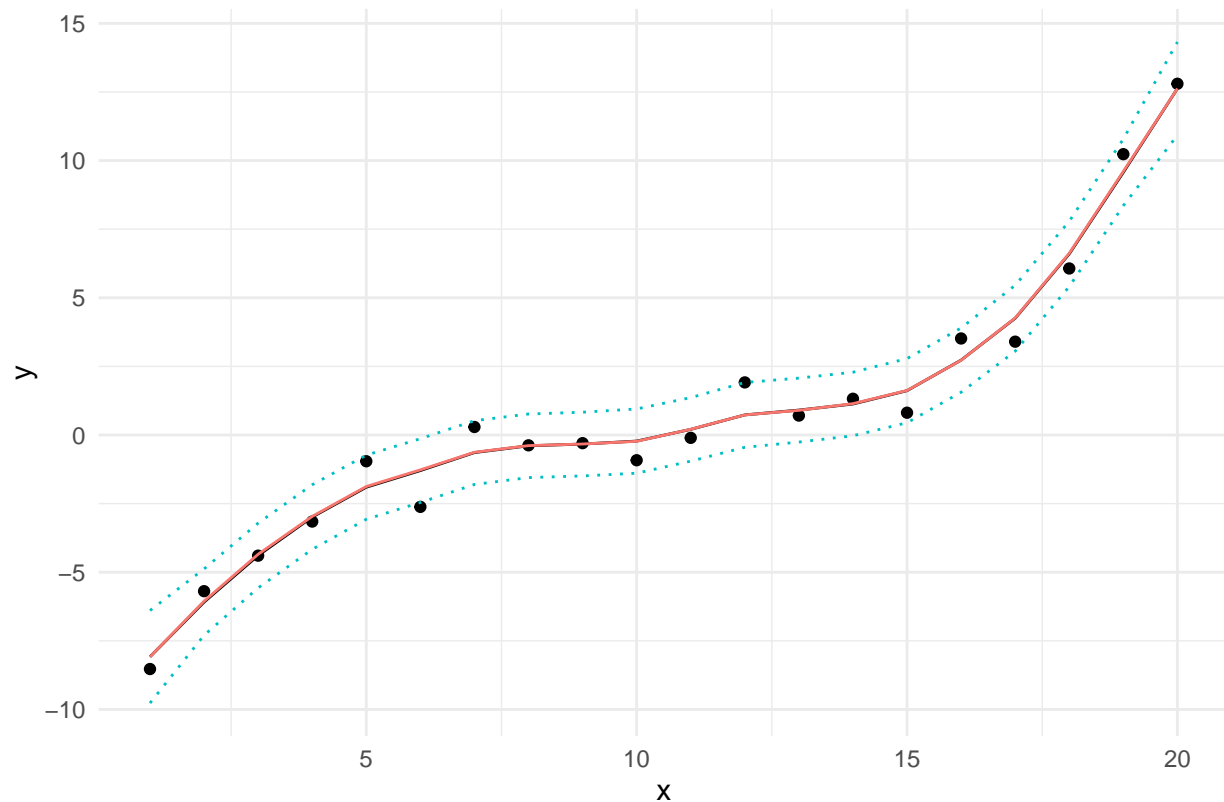
Comparison with results obtained in part 2:

```
#Extracting mean and variance from the INLA model.
mean <- fit$summary.random$x$mean
sd <- fit$summary.random$x$sd
lower <- mean - qnorm(0.025, lower.tail = FALSE)*sd
upper <- mean + qnorm(0.025, lower.tail = FALSE)*sd

inladf <- cbind(x, mean, lower, upper)
inladf <- as.data.frame(inladf)

#GGplot
plot <- ggplot(data = inladf, aes(x = x, y = mean)) + geom_point(aes(x = x, y = gaus$V1 )) + geom_line(aes(x = x, y = mean))
```

```
plot <- plot + labs(y = "y", x = "x", caption = "Estimate of the smooth effect, where the dotted lines  
plot + theme(legend.position="none") + geom_line(aes(x = x, y = meanvec, color = "blue"))
```

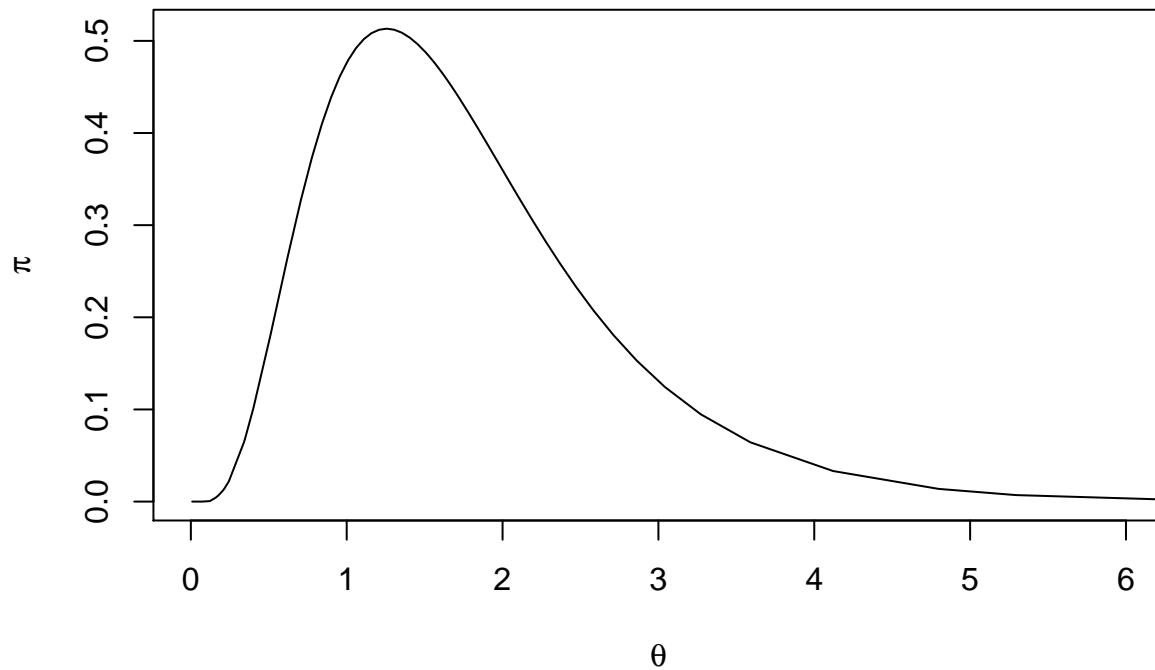


Estimate of the smooth effect, where the dotted lines represents a 95% confidence interval

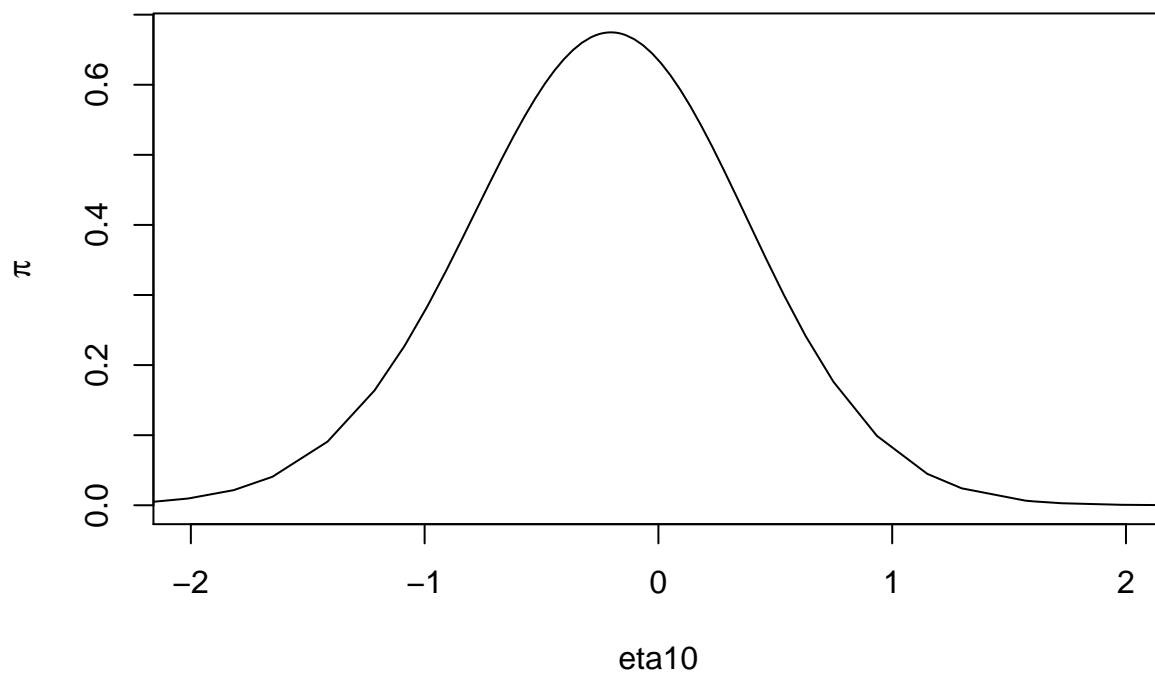
The graph compares the INLA method and the one obtained by MCMC in task 2. We see that the lines that represents the mean are overlapping to a degree that we can barely see that there are two different lines.

**Comparison with results obtained in part 3 and 4:**

```
#Extracting theta and eta from the model and make simple plots.  
theta3 <- fit$marginals.hyperpar$`Precision for x`  
eta3 <- fit$marginals.random$x$index.10  
  
plot(theta3, xlim = c(0,6), type = "l", xlab = bquote(theta), ylab = bquote(pi))
```



```
plot(eta3, xlim = c(-2,2), type = "l", xlab = bquote(eta10), ylab = bquote(pi))
```



The first plot shows the the values of  $\pi(\theta | y)$  for values of theta from 0-6 with the built in INLA scheme. The second plot shows the marginal posterior distribution for the smooth effect for  $\eta_{10}$  with the built in INLA scheme. Wee see that both the plots looks alike to the corresponding plots obtained in the previous tasks.